


Daan Nilis, Julien Lamour, Linus
Handschin

Fast GP-UCB

(Gaussian Process Upper Confidence Bound)


Inspired by:

Srinivas, N., Krause, A., Kakade, S. M., & Seeger, M.
(2009). Gaussian process optimization in the bandit
setting: No regret and experimental design.





GPUCB - Algorithm



- **GOAL:** Find maximum of an expensive function (f)
example: highest temperature in building by turning on sensors
- **INPUT:** search grid, sampling function, #iterations (I)
- **OUTPUT:** mean and variance for every grid point (N^2)
- **RUNTIME:** $O(N^2 * I^3)$
 - I : iterations (how often to sample f)
 - N : search grid size (2D)





Example: $l=0$





Example: $l=1$





Example: $l=2$





Example: $l=3$



Example: $l=4$



Example: $l=5$



Baseline Implementation

- Cost measure: $C(N, \mathbf{I}) = C_{add}(N, \mathbf{I}) + C_{mul}(N, \mathbf{I}) + C_{div}(N, \mathbf{I}) + C_{exp}(N, \mathbf{I})$
- Baseline implementation: linear algebra (argmax + Bayesian update)
- Operational Intensity:


$$W(N, \mathbf{I}) = N^2 * \left(\sum_{i=1}^{\mathbf{I}} (i^2 + i) + \mathbf{I} * k \right) + O(\mathbf{I}^3)$$

$$Q(N) = 4 * N^2$$

$$\begin{aligned} I(N, \mathbf{I}) &= \frac{N^2 * \sum_{i=1}^{\mathbf{I}} (i^2 + i) + O(N^2 \mathbf{I} + \mathbf{I}^3)}{4 * N^2} \\ &= \frac{4I^3}{3} + O\left(\frac{\mathbf{I}^3}{N^2} + \mathbf{I}\right) \gg 1 \end{aligned}$$


Optimization Overview

- 2: $L := \text{cholesky}(K + \sigma_n^2 I)$
 $\alpha := L^\top \setminus (L \setminus \mathbf{y})$
- 4: $f_* := \mathbf{k}_*^\top \alpha$ (mean)
 $\mathbf{v} := L \setminus \mathbf{k}_*$
- 6: $\mathbb{V}[f_*] := k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{v}^\top \mathbf{v}$ (variance)




MMM

- Matrix Matrix Multiplication 8×8
- Matrix Matrix Multiplication $k \times 8$, $k < 8$
- Achieved: 16 f/c




Triangle Solve with kernel computation

- Triangle Solve 8x8, split in:
 - Triangle Solve 4x4
 - MMM 4x4
- Triangle Solve Rest
- Covariance matrix built on the fly




Triangle Solve (with kernel computation on the fly)



gcc 4.8.5 | flags: -O3 -march-avx2 -fno-tree-vectorize | Red Hat 7.3 | search grid 200x200

Fast GPUCB vs Baseline



Fast GPUCB: 30% AVX2 peak
baseline: 12% scalar peak

gcc 4.8.5 | flags: -O3 -march-avx2 -fno-tree-vectorize | Red Hat 7.3 | search grid 200x200